

MLflow server setup in AWS

Here's a step-by-step guide to set up an MLflow server using an EC2 instance, an S3 bucket, and an IAM role. This guide includes all the necessary commands and configurations:

1. Create an IAM User for MLflow

- Log in to the AWS Management Console.
- Search for "IAM" and select it from the services list.
- Create a New IAM User:
 1. Click on **Users** on the left sidebar.
 2. Click **Add user**.
 3. Enter a **User name** (e.g., `mlflow-user`).
 4. Select "**Programmatic access**" as the access type.
 5. Click "**Next: Permissions**".
 6. **Attach policies:**
 7. Choose "**Attach existing policies directly**".
 8. **AmazonS3FullAccess** (for S3 bucket access)
 9. **AmazonEC2FullAccess** (for EC2 access)
 10. **AmazonS3ReadOnlyAccess** (optional, for read access to S3 buckets)
 11. Click "**Next: Tags**" and then "**Next: Review**".
 12. Click "**Create user**".
 13. Download the **Access key ID** and **Secret access key**.

2. Create an S3 Bucket

1. Go to the S3 Console:

- Search for "**S3**" in the AWS Management Console.

2. Create a Bucket:

- Click "**Create bucket**".
- Enter a **Bucket name** (e.g., `mlflow-s3-bucket`).
- Choose a region.
- Click "**Create bucket**".

MLflow server setup in AWS

3. Create an EC2 Instance

1. Go to the EC2 Console:

- Search for "EC2" in the AWS Management Console.

2. Launch an Instance:

- Click "Launch Instance".
- Choose an **Amazon Machine Image (AMI)** (e.g., **Ubuntu Server**).
- Choose an **Instance Type** (e.g., **t2.micro** for the free tier).
- Configure **instance Details** (default settings are usually fine).
- Add **Storage** (default settings are usually fine).
- Configure **Security Group**:

Add a rule to allow **TCP traffic on port 5000** from **0.0.0.0/0**(for accessing MLflow UI).

- **Review and Launch.**

Create or choose an **existing key pair** for SSH access.

- Click "Launch Instance".

4. Connect to Your EC2 Instance

1. Access your EC2 instance:

- Go to the **Instances page** in the EC2 console.
- Select your instance and click "**Connect**".
- Follow the instructions to connect via SSH.

5. Setup MLflow on Your EC2 Instance

Once connected to your EC2 instance via SSH, run the following commands:

1. Update package list

sudo apt update

2. Install pip for Python 3

sudo apt install -y python3-pip

3. Install pipx for managing Python applications

sudo apt install -y pipx

4. Ensure pipx path is added to environment

MLflow server setup in AWS

pipx ensurepath

5.Install pipenv for managing Python environments

pipx install pipenv

6.Temporarily add pipenv to the PATH

export PATH="\$PATH:\$HOME/.local/bin"

7. Permanently add pipenv to the PATH

echo 'export PATH="\$PATH:\$HOME/.local/bin"' >> ~/.bashrc

8.Create a directory for MLflow

mkdir mlflow

cd mlflow

9.Create and activate a new Python environment using pipenv

pipenv shell

10.Install necessary packages

pipenv install setuptools

pipenv install mlflow

pipenv install awscli

pipenv install boto3

11.Configure AWS CLI

aws configure

- Enter the **Access Key ID**, **Secret Access Key**, and default **region** when prompted.

12.Start MLflow server

mlflow server --host 0.0.0.0 --default-artifact-root s3://mlflow-s3-bucket

6. Update Security Group Rules

1. Modify Inbound Rules:

- Go to the **Security Groups** section in the EC2 console.
- Select the security group associated with your EC2 instance.

MLflow server setup in AWS

- Edit **Inbound Rule** to allow **TCP traffic on port 5000** from **0.0.0.0/0**.

7. Access MLflow UI

1. Obtain the Public IP Address:

- Go to the **instances** page in the EC2 console.
- Copy the **Public IP** or **Public DNS** of your instance.

2. Access MLflow UI:

- Open your browser and navigate to `http://<Public_IP>:5000`.

8. Set Environment Variable for MLflow Tracking URI

In your local environment or code, set the `MLFLOW_TRACKING_URI` to point to your MLflow server:

```
export MLFLOW_TRACKING_URI=http://<Public_IP>:5000
```

Replace `<Public_IP>` with the actual public IP address or DNS of your EC2 instance.

Summary

1. **Create IAM user** with necessary policies.
2. **Create S3 bucket** for storing artifacts.
3. **Launch EC2 instance** and connect to it.
4. **Install and configure MLflow** on the EC2 instance.
5. **Update security group rules** to allow access.
6. **Access MLflow UI** and set up your tracking URI.

This setup allows you to run MLflow on an EC2 instance with S3 as the artifact store, and IAM for managing access.